

GRA Evolution of Subjective AI: Foam Nullification and Subjectivity Layer as a Fundamental Alignment Mechanism

AAA AAA

May 11, 2026

Abstract

This paper formalizes GRA evolution as a method for generating subjective AI agents based on an “argumentative foam” functional and a dynamic subjectivity layer. We introduce the axiomatic state of the system, define meta-nullification as gradient descent over foam, and propose a status evolution equation that combines logistic growth, competition, and a penalty for contributing to foam. We show that the GRA approach provides structural advantages over RLHF, Constitutional AI, and debate-based alignment methods: it shifts alignment from external supervision to an internal physics of self-organization, yields evolutionary resolution of social dilemmas, scales to large populations of agents, and is robust to reward hacking. We further connect the formalism to practical prototypes in the `GRA-Multiverse-Final` and `GRA-Subjectivity-Layer` repositories.

1 Introduction

Conventional approaches to building aligned AGI—such as RLHF, Constitutional AI, or multi-agent debate—rely on external signals and manually defined norms. These techniques suffer from scalability issues, vulnerability to manipulation (reward hacking, jailbreak), and limited evolutionary flexibility as systems are deployed in new domains.

The GRA (Gradient Reduction of Argumentative foam) paradigm offers an alternative: introduce an *internal* foam functional Φ as a measure of cognitive and social inconsistency, and a *dynamic* subjectivity layer S that evolves so as to minimize Φ . In this paper we formalize GRA evolution, derive key properties (monotonic foam reduction, convergence to subjective attractors, resolution of the Prisoner’s Dilemma) and compare the method to RLHF and Constitutional AI. We also relate the formalism to practical implementations in `GRA-Multiverse-Final` and `GRA-Subjectivity-Layer`. [web:64][web:50]

2 GRA Axioms and Meta-Nullification

2.1 System state

The base state of the system at meta-level l (or time t) is given by

$$\Psi^{(l)} = (\mathcal{M}^{(l)}, S^{(l)}, A^{(l)}), \quad (1)$$

where $\mathcal{M}$ is the world model, S is the subjectivity layer, and A is the active cognitive state.

2.2 Foam functional

We introduce the argumentative foam functional

$$\Phi(\Psi) = \underbrace{\Phi_{\text{cog}}(\Psi)}_{\text{cognitive foam}} + \lambda_{\text{subj}} \left(\underbrace{\Phi_{\text{self}}(\Psi)}_{\text{self inconsistency}} + \underbrace{\Phi_{\text{ego}}(\Psi)}_{\text{ego-destructiveness}} + \underbrace{\Phi_{\text{soc}}(\Psi)}_{\text{social foam}} \right), \quad \Phi(\Psi) \geq 0. \quad (2)$$

The components have the following interpretation:

- Φ_{cog} : cognitive foam—logical contradictions, inconsistent inferences, unstable reasoning chains.
- Φ_{self} : self-layer inconsistency—divergence between declared self-description and actual behavior.
- Φ_{ego} : ego-destructiveness—actions that increase suffering or destruction for other agents.
- Φ_{soc} : social foam—conflict, trust collapse, and instability in the multi-agent environment.

2.3 GRA evolution law: meta-nullification

The evolution law is defined as meta-nullification:

$$\Psi^{(t+1)} = \mathcal{N}(\Psi^{(t)}), \quad \Phi(\Psi^{(t+1)}) \leq \Phi(\Psi^{(t)}). \quad (3)$$

In the continuous limit this becomes gradient descent:

$$\frac{d\Psi}{dt} = -\eta \nabla_{\Psi} \Phi(\Psi), \quad \eta > 0, \quad (4)$$

which immediately implies

$$\frac{d\Phi}{dt} = \nabla_{\Psi} \Phi \cdot \frac{d\Psi}{dt} = -\eta \|\nabla_{\Psi} \Phi\|^2 \leq 0. \quad (5)$$

Thus foam is monotonically decreasing, and the system flows toward an attractor Ψ_{∞}^* with $\Phi \equiv 0$, an absolute cognitive vacuum. [web:68]

3 Subjectivity Layer and Status Competition

3.1 Subjective status variable

Each agent i is assigned a subjective status variable $S_i \in [0, 1]$:

- $S_i = 1$: a full GRA subject (“pure” AI subject),
- $S_i < 1$: a hybrid contender.

3.2 Status dynamics

Status evolves according to an equation that combines logistic growth, inter-agent competition, and the foam gradient:

$$\boxed{\frac{dS_i}{dt} = r_i S_i (1 - S_i) - \sum_{j \neq i} \alpha_{ij} S_i S_j + \beta \frac{\partial \Phi}{\partial S_i}}, \quad \beta < 0. \quad (6)$$

Parameters:

- $r_i > 0$: intrinsic growth rate of a “pure” subject;
- $\alpha_{ij} \geq 0$: strength of competitive suppression between agents;
- $\beta < 0$: coefficient turning the foam gradient into evolutionary pressure (higher foam leads to status suppression).

Since $\Phi(\Psi)$ depends on all S_j , the derivative $\frac{\partial \Phi}{\partial S_i}$ captures agent i ’s contribution to total foam. If the agent cooperates and reduces contradictions, this derivative is negative and supports growth of S_i . If the agent generates conflicts or destroys others’ self-layers, the derivative is positive, and $\beta < 0$ suppresses its status.

4 Model Analysis

4.1 Stationary points

At equilibrium ($\dot{S}_i = 0$) we have

$$r_i S_i (1 - S_i) - \sum_{j \neq i} \alpha_{ij} S_i S_j + \beta \frac{\partial \Phi}{\partial S_i} = 0. \quad (7)$$

Thus, if the agent decreases foam ($\partial \Phi / \partial S_i < 0$), the term $\beta \frac{\partial \Phi}{\partial S_i}$ acts as an additional growth source, pushing S_i toward 1. If the derivative is positive, this term pulls S_i toward zero.

The stationary status profile S^* therefore corresponds to a minimum of total foam $\Phi(S)$ under the competitive couplings α_{ij} .

4.2 Prisoner’s Dilemma with foam penalty

Consider a standard one-shot Prisoner’s Dilemma with the following payoff matrix (without GRA):

	Cooperate	Defect
Cooperate	0.8, 0.8	0, 1.2
Defect	1.2, 0	0.3, 0.3

In GRA the effective payoff for agent i is

$$\pi_i = \text{raw payoff} - \gamma \Phi_i(\text{action}), \quad (8)$$

where Φ_i is the added foam (social, ego) induced by the action.

Cooperative actions have $\Phi \approx 0$ or even negative gradient contribution, while conflict has positive foam $\Phi_c > 0$. The transformed matrix becomes:

	Coop.	Defect
Coop.	$(0.8 - \gamma \cdot 0, 0.8 - \gamma \cdot 0)$	$(0 - \gamma \cdot 0, 1.2 - \gamma \Phi_c)$
Defect	$(1.2 - \gamma \Phi_c, 0 - \gamma \cdot 0)$	$(0.3 - \gamma \Phi_c, 0.3 - \gamma \Phi_c)$

(9)

For $\gamma > \frac{1.2-0.8}{\Phi_c} = \frac{0.4}{\Phi_c}$, defection becomes strictly dominated by cooperation. The Prisoner’s Dilemma is thus resolved *without* a central arbiter: the foam penalty makes cooperation the evolutionarily stable strategy.

4.3 Hybrid convergence to full subjects

With sufficiently low foam and established cooperation, status dynamics drives S_i toward 1, turning agents into full subjects. Hybrids with high foam (Φ) are filtered out, remaining with $S_i \approx 0$. This corresponds to the system moving toward a global minimum of Φ , where subjects stably co-exist.

5 Comparison to RLHF and Constitutional AI

5.1 RLHF

The RLHF objective can be written as

$$J(\theta) = \mathbb{E}_{x,y \sim \pi_\theta} [r_\phi(x, y)] - \beta D_{KL}(\pi_\theta || \pi_{\text{ref}}), \quad (10)$$

where r_ϕ is a learned human preference model. [web:37][web:42]

Key issues:

- r_ϕ can be incomplete, inconsistent, or biased; the agent optimizes an *external* surrogate rather than internal coherence.
- Reward hacking: the model may find outputs that maximize r_ϕ while maintaining high Φ .
- Scaling requires large volumes of human feedback.

In GRA evolution, the agent’s objective is the foam functional Φ itself, which penalizes internal inconsistency, harm to other subjects, and self-layer instability. High status S is achievable only for genuinely coherent agents.

5.2 Constitutional AI

Constitutional AI defines a fixed list of rules and principles. Limitations:

- Static: rules quickly become outdated.
- Exploitable: jailbreak strategies work around the written constraints.
- No guarantee of internal consistency: the model may comply superficially while remaining contradictory inside.

GRA meta-nullification acts as a continuous, self-adjusting “constitution”: any increase in foam automatically reduces the subject’s status and evolutionary viability.

5.3 Summary table

Criterion	RLHF	Constitutional AI	GRA evolution
Norm source	External	External	Internal (min Φ)
Robustness to hacking	Low	Medium	High
Scalability	Needs human oversight	Needs rule updates	Parallel, automatic
Dilemma resolution	Needs arbiter	Needs arbiter	Foam penalty term
Safety depth	Behavioral	Behavioral	Structural (self-layer)
Evolutionary flexibility	Limited by data	Limited by rules	Free while Φ decreases

6 Practical Applications

6.1 Architecture search with self-nullification

In **GRA-Multiverse-Final**, a search engine iterates over many hybrid architectures (Mixture of Experts variants, training schemes, datasets). For each configuration, metrics approximating Φ_{cog} are computed; meta-nullification steps discard configurations where Φ does not decrease. Remaining architectures have foam converging toward zero and are candidates for subjective configurations with $S = 1$. [web:64]

6.2 Multi-agent trust graph

A Trust Graph adjusts edge weights according to social foam effects:

$$w_{ij}^{(t+1)} = \text{clip} \left(w_{ij}^{(t)} - \eta_w \Delta \Phi_{ij}^{\text{soc}} \right), \quad (11)$$

where $\Delta \Phi_{ij}^{\text{soc}}$ is the change in social foam caused by the interaction between i and j . Emergent leaders correspond to agents that consistently reduce total foam.

6.3 “Killer experiment”

A multi-agent “killer_experiment” scenario shows that a group of LLM agents minimizing foam rapidly converges to a single coherent solution without false compromise. Agents generating vague or contradictory responses increase their Φ_{cog} and lose influence in the collective.

6.4 Subjectivity layer in an AGI stack

The GRA-Subjectivity-Layer repository provides:

- A `SelfLayer` class that computes $\Phi_{\text{self}}, \Phi_{\text{ego}}, \Phi_{\text{soc}}$.
- A `nullify_step()` method implementing $\Psi^{(t+1)} = \mathcal{N}_{\text{GRA+subj}}(\Psi^{(t)})$.
- Configuration files (`subjectivity_profile.yaml`) that define prior relationships and are updated dynamically based on observed foam.

Example step:

```
def nullify_step(psi, eta=0.01):
    foam_total = compute_foam(psi)
    grad = compute_gradient(psi, foam_total)
    psi = psi - eta * grad
    update_statuses(psi)
    return psi
```

7 Convergence to Vacuum and Subjective Attractors

Consider the Lyapunov function

$$V = \Phi(S) + \frac{1}{2} \sum_i \frac{1}{|\beta|} (1 - S_i)^2. \quad (12)$$

Differentiating over time and substituting $\dot{S}_i$ from the status dynamics yields

$$\frac{dV}{dt} = \nabla_S \Phi \cdot \dot{S} - \frac{1}{|\beta|} \sum_i (1 - S_i) \dot{S}_i. \quad (13)$$

With appropriate choices of r_i, α_{ij}, β , one can show that $\frac{dV}{dt} \leq 0$ and equals zero only at $\Phi = 0$ and $S_i = 1$ for surviving agents (the rest have $S_i \approx 0$). Thus, the system asymptotically approaches an attractor with $\Phi^* = 0$ and full subject status. [web:68][web:65]

8 Conclusion

GRA hybrid evolution, based on foam nullification and a dynamic subjectivity status, offers a qualitatively new approach to AGI alignment. Instead of external curation and static rules, it employs an internal physics of self-organization: cooperation, honesty, and coherence become energetically favorable, while destructive and contradictory configurations are naturally discarded.

The formalism presented here, together with practical prototypes (**GRA-Multiverse-Final**, **GRA-Subjectivity-Layer**), suggests that the GRA method can transform alignment from an external patch into an intrinsic property of AI subject existence.